

PS3 Review Session

Jingwei Ji


Overview

Space carving

Single Object Recognition via SIFT

Histogram of Oriented Gradient (HOG)

Space carving - overview




Visual hull:
an upper bound estimate

Space carving - overview

Steps:


- Estimate silhouettes of images (could be based on some heuristics, e.g. color)
- Form the initial voxels as a cuboid
- Iterate over cameras and remove the voxels which project to the dark part of each silhouette



Space carving - (a) (b) (c)

Steps:


- Estimate silhouettes of images (could be based on some heuristics, e.g. color)
- Form the initial voxels as a cuboid #TODO
- Iterate over cameras and remove the voxels which project to the dark part of each silhouette #TODO



Space carving - (a) (b) (c)

Steps:


- Estimate silhouettes of images (could be based on some heuristics, e.g. color)
- Form the initial voxels as a cuboid
 - Try not just do for loops, too slow
 - Probably useful functions: `np.meshgrid`, `np.repeat`, `np.tile`
- Iterate over cameras and remove the voxels which project to the dark part of each silhouette
 - Critical step: project 3D voxel into 2D
 - Remember using homogenous coordinates for projection
 - Also, better no for loops...




Space carving - (d)

Steps:


- Estimate silhouettes of images (could be based on some heuristics, e.g. color)
- Form the initial voxels as a cuboid
 - Improvement: tighter bounded cuboid
 - How: do a coarse carving first (we use num_voxels = 4000)
- Iterate over cameras and remove the voxels which project to the dark part of each silhouette



Coarse
Carving



Final Output



Space carving - (e)


Steps:

- Estimate silhouettes of images (could be based on some heuristics, e.g. color)
 - Issues: the quality of silhouettes
 - The silhouette from each camera is not perfect, but the result is ok. Why?
 - Experiment: use only a few of the silhouettes
- Form the initial voxels as a cuboid
- Iterate over cameras and remove the voxels which project to the dark part of each silhouette

SIFT descriptor

David G. Lowe. ["Distinctive image features from scale-invariant keypoints."](#) IJCV 60 (2), 04

- Alternative representation for image regions
- Location and characteristic scale s given by DoG detector



- 1 Compute gradient at each pixel
- 2 $N \times N$ spatial bins
- 3 Compute an histogram h_i of M orientations for each bin i
- 4 Concatenate h_i for $i=1$ to N^2 to form a $1 \times MN^2$ vector H
- 5 Gaussian center-weighting
- 6 Normalize to unit norm


Typically $M = 8$; $N = 4$
 $H = 1 \times 128$ descriptor

Single Object Recognition Via SIFT - (a)

We've implemented SIFT descriptor for you, and your task is using it for object recognition.

Read: **Section 7. Application to object recognition** in Lowe's SIFT paper
<http://www.cs.ubc.ca/~lowe/papers/ijcv04.pdf>

Be sure to understand why is there a threshold and how to use it.




Single Object Recognition Via SIFT - (b)

RANSAC to refine matching

Basic idea: use RANSAC to fit a homography matrix H between two set of key points. Only keep the inliers for matching, remove the outliers.

Think: how many pairs of correspondence do we need in each iteration?



Should be removed

More about RANSAC

Theoretical properties: see lecture notes / slides, figuring out the relations between

- N : number of samples
- e : outlier ratio
- s : minimum number of data points to fit the model
- p : probability guaranteed

Single Object Recognition Via SIFT - (d) (e)

Find quantitative relation between two bounding boxes

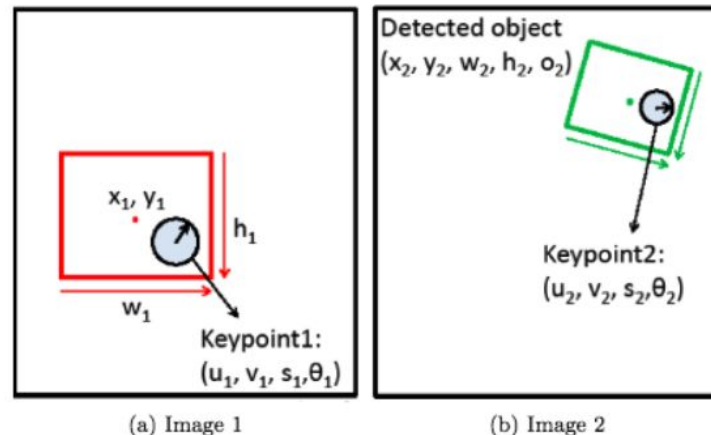
Input (u_1, v_1, s_1, θ_1 ; u_2, v_2, s_2, θ_2 ; x_1, y_1, w_1, h_1), find (x_2, y_2, w_2, h_2, o_2)

Keypoint1

Keypoint2

Bbx1


Bbx2 & Orientation



HoG = Histogram of Oriented Gradients

Navneet Dalal and Bill Triggs, Histograms of Oriented Gradients for Human Detection, CVPR05

- Like SIFT, but...
 - Sampled on a dense, regular grid around the object
 - Gradients are contrast normalized in overlapping blocks



Histogram of Oriented Gradients (HOG) - Overview

Three levels of abstraction:

1. pixel-wise gradients (angle, magnitude) -> histogram in cell
2. histogram in cell -> HoG feature in the block
3. HoG feature in the block -> total HoG feature in the image

See whiteboard

Histogram of Oriented Gradients (HOG) - Overview

- Compute the gradients of image (angles, magnitudes)
- Divide the image into blocks with overlapping, each image contains many blocks
- Divide blocks into cells, each block contains several cells, and each cell contains several pixels.
- Generate the histogram (of shape (nbins,)) of each cell, the concatenation of histograms of cells in a block serves as the feature vector of this block
- Output feature: (H_blocks, W_blocks, cells_in_block ** 2 * nbins)

Histogram of Oriented Gradients (HOG) - (a)

See whiteboard

- Compute the gradients of image (angles, magnitudes)
- Divide the image into blocks with overlapping, each image contains many blocks
- Divide blocks into cells, each block contains several cells, and each cell contains several pixels.
- Generate the histogram (of shape (nbins,)) of each cell, the concatenation of histograms of cells in a block serves as the feature vector of this block
- Output feature: $(H_blocks, W_blocks, cells_in_block ** 2 * nbins)$

Histogram of Oriented Gradients (HOG) - (a)

Compute the gradients of image (angles, magnitudes)

The way that the angles and magnitude per pixel are computed as follows:
Given the following pixel grid

```
P1 P2 P3
P4 P5 P6
P7 P8 P9
```

We compute the angle on P5 as $\arctan(dy/dx) = \arctan(P2-P8 / P4-P6)$.

The magnitude is simply $\sqrt{(P4-P6)^2 + (P2-P8)^2}$

Histogram of Oriented Gradients (HOG) - (b)

- Compute the gradients of image (angles, magnitudes)
- Divide the image into blocks with overlapping, each image contains many blocks
- Divide blocks into cells, each block contains several cells, and each cell contains several pixels.
- **Generate the histogram (of shape (nbins,)) of each cell**, the concatenation of histograms of cells in a block serves as the feature vector of this block
- Output feature: (H_blocks, W_blocks, cells_in_block ** 2 * nbins)

Histogram of Oriented Gradients (HOG) - (b)

```
def generate_histogram(angles, magnitudes, nbins = 9):
```

Binning the angles with magnitudes as weights

Be careful about the edge cases: what if the angle is close to 0 or 180 degrees?

Histogram of Oriented Gradients (HOG) - (c)

- Compute the gradients of image (angles, magnitudes)
- Divide the image into blocks with overlapping, each image contains many blocks
- Divide blocks into cells, each block contains several cells, and each cell contains several pixels.
- Generate the histogram (of shape (nbins,)) of each cell, the concatenation of histograms of cells in a block serves as the feature vector of this block
- Output feature: (H_blocks, W_blocks, cells_in_block ** 2 * nbins)

