

CS231A Midterm Review

Mar 2, 2018

Midterm Logistics

- In-class midterm at Skilling Auditorium at 1:30-2:50 PM on Monday.
- SCPD students not taking exam at Stanford should coordinate with SCPD.
Let us know if you are coming to stanford so we can accommodate.
- Open book and open notes. Not open computer.
- Lectures 1 - 12 (through Image Classification & 2D Object Detection)
- 10 T/F, 10 MC, 5 short questions, 4 longer questions.
- Bring a calculator to solve numerical questions.

Topics you should know for the Exam


- General knowledge of linear algebra (matrix multiplication, SVD, etc)
- Camera Models and Transformations
- Non-perspective Cameras
- Camera Calibration
- Single View Metrology
- Epipolar Geometry
- Structure from Motion
- Active Stereo and Volumetric Stereo
- Fitting and Matching
- RANSAC
- Hough Transform
- Detectors and Descriptors
- Image Classification
- 2D Object Detection

Homogeneous Coordinates

- Homogeneous coordinates allow us to apply a larger variety of transformations with matrix multiplication
 - For example, we use homogeneous coordinates to handle the 3D \rightarrow 2D projection
- Any point (x, y) becomes represented as $(x, y, 1)$
- More generally $(a_1, a_2, \dots, a_n, w)$ represents the point $(a_1/w, a_2/w, \dots, a_n/w)$

Types of transformations

- Isometric transformations preserve distances
 - Rotation, translation, reflection
- Similarity transformations preserve shape
 - Rotation, translation, scaling
- Affine transformations preserve parallelism
 - Rotation, translation, scaling, shearing, etc.
 - $T(v) = Av + t$, A is invertible
- Projective transformations map lines to lines
 - Pretty much everything else



Examples in class are 2D, need to know how to generalize to 3D

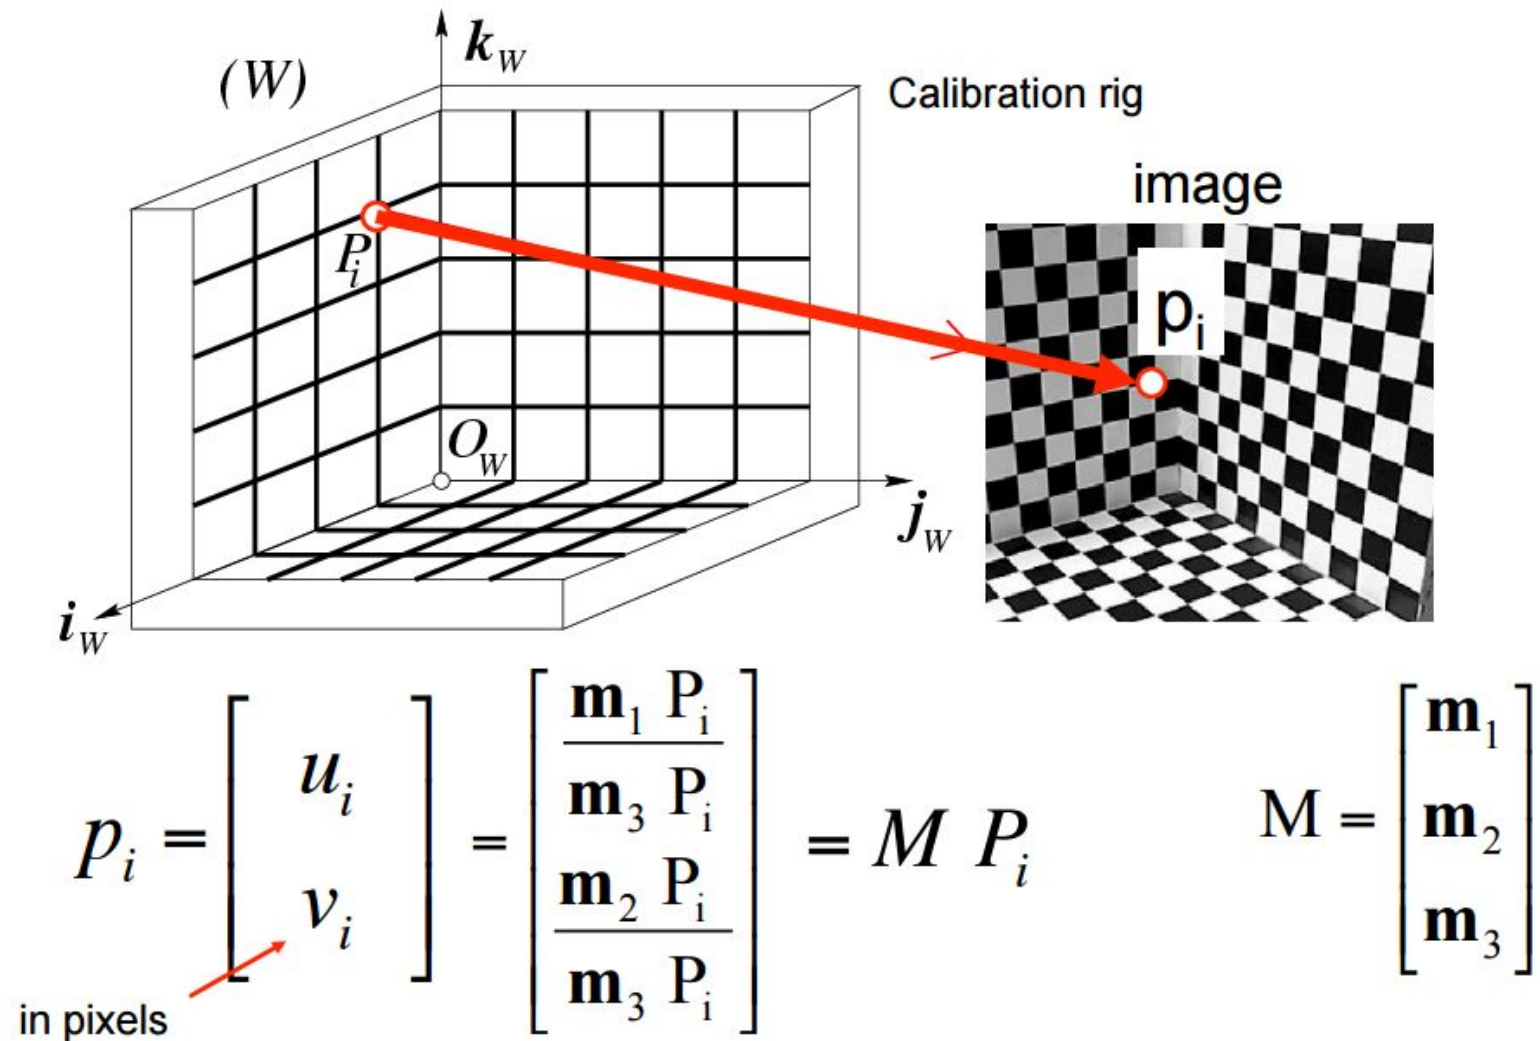
Camera Parameters

- Extrinsic parameters
 - Rotation and translation from the world frame
- Intrinsic parameters
 - Focal length in x and y direction, camera center offset, skew, distortion
 - Most people assume only 5 parameters (for the sake of this class)

$$K = \begin{bmatrix} \alpha & -\alpha \cot \theta & c_x \\ 0 & \frac{\beta}{\sin \theta} & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

$$P' = K \begin{bmatrix} R & T \end{bmatrix} P_w = MP_w$$

Camera Calibration



Camera Calibration

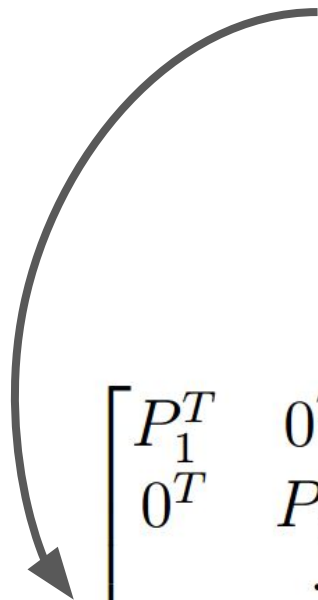
$$u_1(m_3 P_1) - m_1 P_1 = 0$$

$$v_1(m_3 P_1) - m_2 P_1 = 0$$

$$\vdots$$

$$u_n(m_3 P_n) - m_1 P_n = 0$$

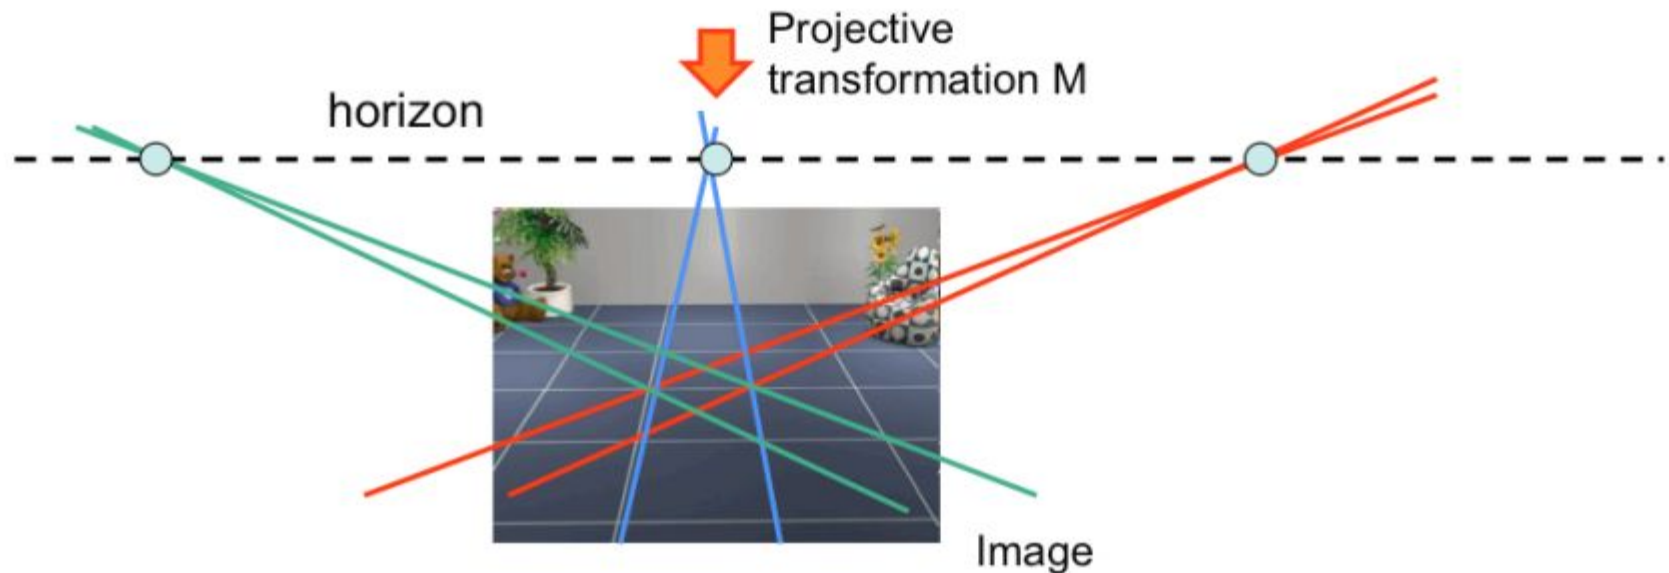
$$v_n(m_3 P_n) - m_2 P_n = 0$$


$$\begin{bmatrix} P_1^T & 0^T & -u_1 P_1^T \\ 0^T & P_1^T & -v_1 P_1^T \\ \vdots & \vdots & \vdots \\ P_n^T & 0^T & -u_n P_n^T \\ 0^T & P_n^T & -v_n P_n^T \end{bmatrix} \begin{bmatrix} m_1^T \\ m_2^T \\ m_3^T \end{bmatrix} = \mathbf{P} \mathbf{m} = 0$$

Solve by SVD!

Single View Metrology

- Vanishing points and vanishing lines (horizon)




- This leads to being able to find angles between lines and planes (recall PS1)
- You can also calibrate the camera from a single image!

$$\cos \theta = \frac{\mathbf{v}_1^T \boldsymbol{\omega} \mathbf{v}_2}{\sqrt{\mathbf{v}_1^T \boldsymbol{\omega} \mathbf{v}_1} \sqrt{\mathbf{v}_2^T \boldsymbol{\omega} \mathbf{v}_2}}$$

[Eq. 28]


$$\boldsymbol{\omega} = (\mathbf{K} \mathbf{K}^T)^{-1}$$

Epipolar Geometry



- Understanding the geometry of the scene and the cameras
- Should have knowledge of this entire scene and basic triangulation
 - Epipoles, epipolar lines, reprojection error, etc.

Unique Cases of Epipolar Geometry



- Parallel cameras make the epipoles at infinity
- Forward translation make the epipoles in the same location

The Fundamental Matrix

$$\mathbf{p}^T \mathbf{F} \mathbf{p}' = 0$$


$$\mathbf{F} = \mathbf{K}^{-T} \cdot [\mathbf{T}_\times] \cdot \mathbf{R} \mathbf{K}'^{-1}$$

- Relates corresponding points with a single constraint
- 7 degrees of freedom
- Can be found using Eight Point Algorithm and Normalized Eight-Point algorithm

$$\begin{bmatrix} u_1 u'_1 & v_1 u'_1 & u'_1 & u_1 v'_1 & v_1 v'_1 & v'_1 & u_1 & v_1 & 1 \\ u_2 u'_2 & v_2 u'_2 & u'_2 & u_2 v'_2 & v_2 v'_2 & v'_2 & u_2 & v_2 & 1 \\ u_3 u'_3 & v_3 u'_3 & u'_3 & u_3 v'_3 & v_3 v'_3 & v'_3 & u_3 & v_3 & 1 \\ u_4 u'_4 & v_4 u'_4 & u'_4 & u_4 v'_4 & v_4 v'_4 & v'_4 & u_4 & v_4 & 1 \\ u_5 u'_5 & v_5 u'_5 & u'_5 & u_5 v'_5 & v_5 v'_5 & v'_5 & u_5 & v_5 & 1 \\ u_6 u'_6 & v_6 u'_6 & u'_6 & u_6 v'_6 & v_6 v'_6 & v'_6 & u_6 & v_6 & 1 \\ u_7 u'_7 & v_7 u'_7 & u'_7 & u_7 v'_7 & v_7 v'_7 & v'_7 & u_7 & v_7 & 1 \\ u_8 u'_8 & v_8 u'_8 & u'_8 & u_8 v'_8 & v_8 v'_8 & v'_8 & u_8 & v_8 & 1 \end{bmatrix} \begin{bmatrix} F_{11} \\ F_{12} \\ F_{13} \\ F_{21} \\ F_{22} \\ F_{23} \\ F_{31} \\ F_{32} \\ F_{33} \end{bmatrix} = 0$$

Solve by SVD!




Structure from Motion




- Estimating both the camera positions and the 3D structure simultaneously from point correspondences
- You've implemented a few algorithms:
 - Factorization method
 - An iterative triangulation method

Active Stereo and Volumetric Stereo

- Active Stereo
 - Replaces one camera with a projector
- Volumetric Stereo
 - Space carving
 - Shadow carving
 - Voxel coloring

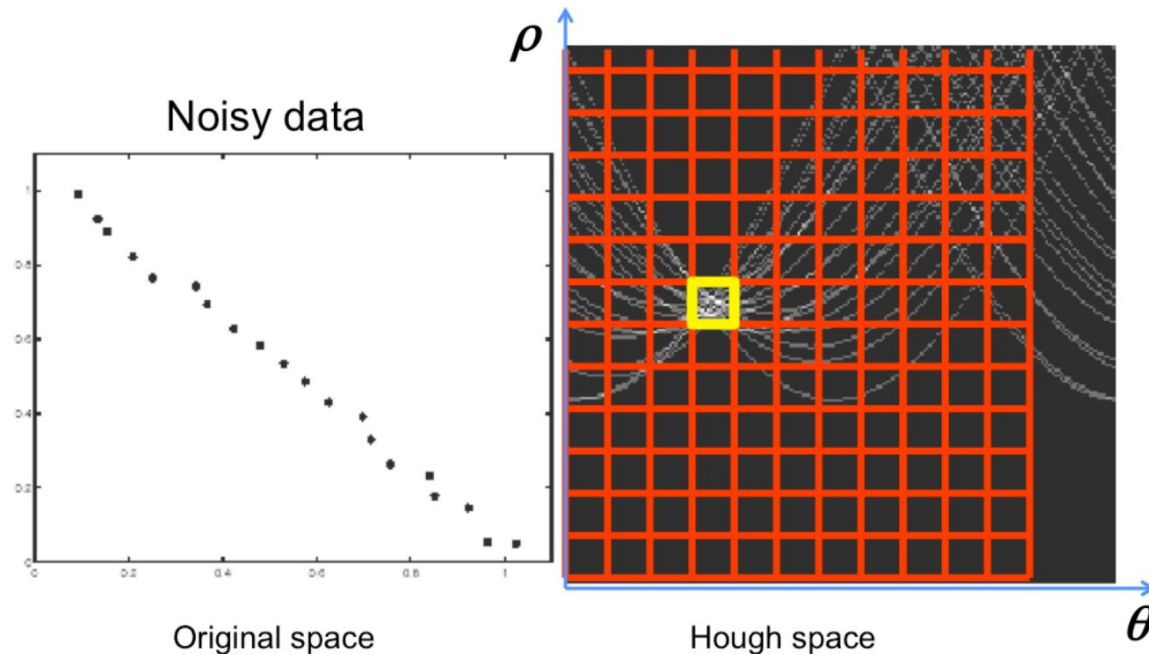


RANSAC



- Select random sample of minimum size
- Compute a model from this
- Compute the inliers within the model
- Repeat steps for a fixed amount and return the model with the most inliers

Hough Transform



- Find some parameter space that defines the line, plane, etc. that we're trying to estimate
- For each observation, plot in this parameter space
 - Could be points, lines, hyperplanes, etc.
- Grid up the parameter space and find cells with many observations

Detectors and Descriptors

- Corner detectors
 - Harris corner detector
- Edge detectors
 - Find areas of high gradients, but should smooth before doing so to remove noise
 - Should know about Laplacian of Gaussian and Difference of Gaussian
- Blob detection
 - Similar to edge detection, but in 2D
- SIFT
 - A local descriptor around keypoints based on gradients in the image
 - Scale and in-plane rotation invariant
 - Steps to calculate SIFT descriptor
- HOG
 - Implemented in PS3 - you should know about it!
 - Steps to calculate HOG feature

Image Classification and 2D Object Detection

- Bag of words
 - Histogram representation of “words” (features)
- Part of PS4 (don’t need to implement, but skimming the ideas in preparation for the midterm is a good idea)
 - Sliding window detectors
 - Non-maximal suppression

Exam Advice

- When studying, make a cheat sheet to quickly reference at exam time.
 - Since you have 80 minutes, you do not want to sift through pages of notes.
- You will not need to know the complex math derivations involved in the course.
 - There are some linear algebra problems, though!
- You will need to know how generally things work and explain them (camera matrices, SFM, RANSAC, Hough transforms, etc.).
- This review session is not comprehensive of all material on the exam!